

ABHIRAJ KARPE

Pune, Maharashtra, India | abhirajkarpe7@gmail.com | +91 9975956475

[linkedin.com/in/abhiraj-karpe](https://www.linkedin.com/in/abhiraj-karpe) | github.com/abhh10

PROFESSIONAL SUMMARY

Data Engineer with hands-on experience building batch, streaming, and CDC pipelines using Apache Spark, Kafka, Databricks, Airflow, dbt, and Snowflake. Experienced in incremental processing, data modeling, lakehouse architecture, and pipeline optimization, with production engineering experience at Jio Platforms. Seeking to build scalable and reliable data platforms.

EXPERIENCE

Software Engineering Intern — Jio Platforms Limited

Jan 2026 – Jul 2026

Pune, Maharashtra, India

- Converted legacy ADB script outputs from raw text to structured JSON across OTT device/app profiling workflows, reducing storage overhead by **30–40%** and cutting downstream parsing time by an estimated **~35%**.
- Optimized memory and CPU profiling workflows for OTT applications and devices, reducing job runtime by an estimated **~20%** and improving data collection consistency across device models.
- Designed and deployed a production ingestion API for crash and profiling data, integrating ADB tooling, custom parsing logic, and database storage into a pipeline used across daily QA test cycles.
- Built automated collection of device-level metrics (**RAM usage, foreground apps, storage stats**) using Python and shell scripts, emitting structured JSON for backend ingestion.
- Debugged data-loss and parsing issues across log parsing, ADB, JSON/TXT handling, and backend database storage, ensuring accurate logging of key fields and improving pipeline reliability.

PROJECTS

CDC Lakehouse Pipeline

github.com/abhh10/cdc-lakehouse

PostgreSQL, Debezium, Apache Kafka, Apache Iceberg, MinIO, Hive Metastore, Trino

- Built an end-to-end CDC pipeline streaming PostgreSQL inserts, updates, and deletes through **Debezium and Kafka** into Apache Iceberg tables with **1.03s P95** latency.
- Integrated **MinIO, Hive Metastore, and Trino** for queryable lakehouse analytics; containerized a 7-service stack for reproducible deployment.

Databricks Taxi Analytics Pipeline

github.com/abhh10/databricks-nyc-trip-pipeline

Databricks, PySpark, Delta Lake, Unity Catalog, AWS S3

- Built an incremental Bronze–Silver–Gold pipeline processing **355M+** taxi trip records with schema normalization, data quality filtering, and deduplication.
- Designed partition-aware incremental queries that cut data scanned by **89%** and execution time by **33%**, producing 3 analytics-ready Gold datasets.

Real-Time Crypto Data Engineering Pipeline

github.com/abhh10/crypto-data-engineering-pipeline

Apache Kafka, AWS S3, Snowpipe, Snowflake, dbt, Airflow

- Built a real-time pipeline streaming **CoinGecko** market events through Kafka into Amazon S3 and **Snowflake** via **Snowpipe**.
- Developed **incremental dbt** transformations with MERGE-based upserts across RAW, STAGING, and GOLD layers, orchestrated with **Airflow**.

TECHNICAL SKILLS

Languages: Python, SQL, Java

Data Engineering: Apache Spark, PySpark, Kafka, Debezium, Apache Iceberg, Airflow, dbt, CDC, ETL/ELT, Data Modeling, Incremental Processing, Medallion Architecture

Databricks: Delta Lake, Unity Catalog, Lakebase, Lakeflow Declarative Pipelines, Auto Loader, Databricks Workflows

AWS: S3, Lambda, Glue, Athena, IAM, CloudWatch, Boto3, AWS CLI

Warehousing & Databases: Snowflake, Snowpipe, Amazon Redshift, BigQuery, PostgreSQL, MySQL, MongoDB

AI & LLM: LLM-Assisted Development, Prompt Engineering, RAG, OpenAI API, LangChain, Vector Databases (ChromaDB)

Tools & Backend: Git, GitHub, CI/CD, Docker, Linux, MinIO, Hive Metastore, REST APIs, Node.js, Postman

EDUCATION

Ajeenkya D.Y. Patil University — B.Tech, Computer Science (Software Engineering)

Sept 2022 – Jun 2026

Pune, Maharashtra, India | CGPA: 8.86

CERTIFICATIONS & ACHIEVEMENTS

- AWS re/Start Graduate — Amazon Web Services Training and Certification
- Introduction to Data Engineering — DeepLearning.AI & Amazon Web Services (Coursera)
- Solved 700+ problems on LeetCode; top 0.5% of users in 2025 (Contest Rating: 1575). GeeksforGeeks Coding Score: 313.